

OMol25 Paradox v5h

The OMol25 Paradox

Training-Data Domain Specificity Outweighs Architecture in Neural Network Potential Performance on Zn^{2+} Metalloenzyme Conformer Energies

Abstract

Machine-learning interatomic potentials (MLIPs) trained at high-level DFT are widely assumed to dominate semi-empirical and materials-trained alternatives on molecular tasks. We test this assumption on the conformational energy landscape of fifteen Zn^{2+} -binding matrix metalloproteinase-1 (MMP-1) inhibitors drawn from ChEMBL, generated by Boltz-2x cofold sampling with physicality-steering (--use_potentials). Across thirteen independent Boltz-2x diffusion replicates (Boltz seeds 42–54, conditions v11–v23), we evaluate ten energy functions on 100 conformers per ligand per replicate (19,500 conformers per replicate; 253,500 conformer single points total): three semi-empirical (GFN1-xTB, GFN2-xTB, GFN-FF), one Acellera-trained NN force field (AceFF-2), two molecular NN potentials (AIMNet2-NSE, ANI-2x), three universal MLIPs trained on materials physics (Orb-v3 OMat, Orb-v3 OMol25, MatterSim-5M), and two classical force fields (MMFF94, UFF). Per-ligand intra-conformer Pearson correlations are computed within each ligand (size-invariant by construction), then averaged across the 15 ligands and 13 replicates. Three clusters emerge with high cross-replicate stability ($\text{SD} \leq 0.038$). Cluster I (QM-like) tightly groups GFN1, GFN2, AIMNet2, AceFF-2, Orb-v3 OMat, and MatterSim-5M (intra-cluster $r = 0.66$ – 0.99). Cluster II (force-field) groups MMFF94 and UFF ($r = 0.66$), weakly anti-correlated with Cluster I ($r = -0.13$ to -0.25). The headline finding is the **OMol25 paradox**: the same Orb-v3 architecture trained on Meta FAIR’s OMol25 molecular dataset ($\omega\text{B97M-V/def2-TZVPD}$) drops to $r = 0.374 \pm 0.025$ against GFN1-xTB, versus $r = 0.886 \pm 0.009$ for Orb-v3 OMat — a 0.512 Pearson gap with $\approx 68\sigma$ statistical significance across 13 replicates. The paradox is corroborated by a second materials-trained MLIP, MatterSim-5M, which agrees with Orb-v3 OMat at $r = 0.936 \pm 0.006$ (the tightest NN–NN pair in our matrix) but drops to $r = 0.410 \pm 0.026$ against Orb-v3 OMol25. Cluster placement is set by training-data domain, not by architecture. We discuss the mechanistic explanation — OMol25’s self-reported weakness on “long-range interactions” (Levine et al., arXiv:2505.08762) coincides with the dominant physics of MMP-1 hydroxamate– Zn^{2+} chemistry — and we explicitly retract a preliminary biology-validation claim that was confounded by ligand size. We recommend that practitioners choose materials-trained or broadly-trained MLIPs (Orb-v3 OMat, MatterSim, AIMNet2) over molecular-only-trained MLIPs (Orb-v3 OMol25, SevenNet OMol25) for Zn^{2+} metalloenzyme conformer-ensemble work, pending higher-level DFT validation.

Word count (abstract): 388 words.

1. Introduction

1.1 Matrix metalloproteinase-1 and the Zn^{2+} active site challenge

Matrix metalloproteinase-1 (MMP-1, fibroblast collagenase) is a Zn^{2+} -dependent endopeptidase implicated in tissue remodeling, photoaging, tumor invasion, and fibrotic disease. The catalytic site is a canonical three-histidine zinc triad (HIS-218, HIS-222, HIS-228 in 1HFC numbering, with HID111/HID115/HID121 in our re-numbered system) with a fourth coordination position normally occupied by a hydroxide ion. Clinical-stage MMP-1 inhibitors share a common warhead motif — a hydroxamate ($-\text{CONHOH}$), a thiol, a phosphonate, or a carboxylate — that chelates the catalytic zinc in a bidentate fashion. The remainder of the ligand fills the S1' specificity pocket through a mixture of hydrophobic packing, dipole-dipole interactions, π - π stacking, and dispersion forces. This active-site chemistry is the bane of standard organic-molecule force fields and benchmark-tuned ML potentials, because the relevant physics combines (a) partial charge transfer and back-donation at the Zn^{2+} coordination shell, (b) intramolecular long-range electrostatics shaped by the warhead's strong dipole, and (c) the protein-ligand dispersion that dominates the S1'-pocket interaction.

Consequently, the conformer energy landscape of MMP-1 hydroxamate inhibitors is an unusually demanding probe of any single-point energy method that does not see this physics during training. It is, in particular, a probe of whether modern universal machine-learning interatomic potentials (MLIPs) — trained on millions to billions of DFT calculations spanning either condensed-phase materials chemistry or organic small molecules — generalize transferably to a Zn^{2+} metalloenzyme ligand ensemble.

1.2 The NNP landscape, 2024–2026

A new generation of universal MLIPs has crystallized in the past 18 months. We define four families relevant to our benchmark:

1. **Materials-trained universal MLIPs.** Orb-v3 OMat (Orbital Materials, Apr 2025; Neumann et al., arXiv:2504.10367), MatterSim-5M (Microsoft Research, May 2024; Yang et al., arXiv:2405.04967), and the SevenNet-Omni family (Seoul National University MDIL, Nature Communications 2026) are E(3)-equivariant graph networks trained primarily on the Materials Project, OMat-24, MPA, MatPES, or related condensed-phase DFT datasets. These training sets emphasize periodic crystal structures, broad chemistry (most of the periodic table including transition metals), and configurations sampled across diverse force regimes.
2. **Molecular-trained universal MLIPs.** Orb-v3 OMol25 (Orbital Materials, May 2025) and SevenNet OMol25-high (SNU, 2026) re-use the same Orb-v3 / SevenNet architectures but are trained on Meta FAIR's Open Molecules 2025 dataset (OMol25; Levine et al., arXiv:2505.08762) — 100M+ ω B97M-V/def2-TZVPD single points on 83M small organic molecules. By design, OMol25 prioritizes molecular chemistry accuracy.
3. **Molecular NNPs at smaller scale.** AIMNet2-NSE (Anstine et al., 2024; $\sim 2 \times 10^7$ training points) and ANI-2x (Devereux et al., 2020; $\sim 10^6$ training points) cover H/C/N/O/F/S/Cl-containing organic molecules at modest training scale.

4. **Reactive QM-trained NN force fields.** AceFF-2 (Acellera, 2024) is a special-purpose NN potential trained on protein-ligand QM calculations and intended for drug-discovery workflows.

This generation explicitly claims transferability to molecular tasks. OMol25 in particular has been marketed as “the molecular foundation model” — with the implicit assumption that domain-specialized training (high-level molecular DFT) will outperform broader, materials-leaning training on molecular benchmarks.

1.3 Knowledge gap: cross-NNP rank-order agreement on metalloenzyme ligand ensembles

Most published NNP benchmarks compare predicted versus reference DFT energies on the *training distribution* (e.g., MD17, QM9, SPICE-2, MPF) or on closely-related test sets (PDBbind ligand conformers, ANI-1xnr). What is missing is **cross-NNP rank-order agreement on a hard, well-controlled out-of-distribution task**: how do different MLIP families rank the same set of conformers of the same set of ligands when those ligands chelate a transition metal?

For practitioners, the operational question is: *if I use NNP X for conformer scoring on a Zn^{2+} metalloenzyme target, would NNP Y have given me the same conformer ranking?* If not, the choice of NNP is itself a methodological degree of freedom that must be reported, justified, and ideally cross-validated.

1.4 Our contribution

We address this gap directly. We assemble a panel of ten conformer energy methods spanning the four families above (Section 2.2), generate 13 independent Boltz-2x cofold conformer ensembles of the MMP-1 active site bound to 15 ChEMBL inhibitors (Section 2.1), and compute per-ligand intra-conformer Pearson cross-correlations between every pair of methods (Section 2.3). We deliberately use **per-ligand intra-conformer Pearson** as our primary metric: it is size-invariant by construction (because comparison is within a fixed ligand) and is therefore immune to the ligand-size confounding that we initially overlooked in a related biology-validation analysis (and which we explicitly retract in Section 3.5).

The headline finding — the OMol25 paradox — is that the same Orb-v3 architecture trained on materials versus molecular data lands in two **different** clusters, with the molecular-trained variant agreeing **less** with QM than the materials-trained variant agrees with QM. The paradox is reproduced by an independent architecture (MatterSim, M3GNet-class) trained on materials data, and by an independent architecture (SevenNet-class) trained on the same OMol25 data (paper_A v4). Training-data domain dominates architecture choice for cluster placement.

2. Methods

2.1 Ligand panel and Boltz-2x cofold ensemble generation

Target. Apo monomer of MMP-1 (PDB 1HFC, chain A, catalytic domain, residues 100–262), single Zn^{2+} in the active site, three-histidine first coordination shell. Hydrogens added with `pdb4amber` followed by `tLeap` (AmberTools 23) using the ZAFF (Zinc AMBER Force Field) charge model.

The Zn^{2+} is modeled as a bare 2+ point charge for the purpose of NNP scoring (without ZAFF tetrahedral dummies, which were used only for ABFE preparation in companion paper_B work).

Ligand panel. Fifteen MMP-1 ChEMBL inhibitors selected to span ≈ 4 log units in reported IC50 (4 nM to 18 μM) and four warhead/scaffold classes:

ChEMBL ID	Reported IC50 (nM)	Warhead	Class
CHEMBL443684 (Marimastat)	5	hydroxamate	peptidic
CHEMBL406 (Pri-nomastat)	3	sulfonamide-hydroxamate	non-peptidic
CHEMBL415	4	hydroxamate	non-peptidic
CHEMBL94487	12	hydroxamate	non-peptidic
CHEMBL98	2400	hydroxamate (weak)	aryl-sulfone
CHEMBL57058	—	hydroxamate	non-peptidic
CHEMBL259829 (CGS27023A)	310	non-hydroxamate sulfonamide	non-peptidic
CHEMBL3036	—	hydroxamate	macrocycle
CHEMBL2105729	18,000	hydroxamate (very weak)	non-peptidic
CHEMBL257077	15	hydroxamate	non-peptidic
CHEMBL301236	42	hydroxamate	non-peptidic
CHEMBL292707	200	hydroxamate	non-peptidic
CHEMBL_supp_1 — CHEMBL_supp_3	varied	hydroxamate / carboxylate	supplemental scaffolds

(Heavy-atom counts range from 22 to 74; molecular charges range from -1 to 0 .)

Cofold sampling. Boltz-2x v0.6.x (cofold checkpoint, Mar 2026) is run on the MMP-1 sequence plus each ligand SMILES, with `--use_potentials` enabled (physicality steering active per the protocol established in paper_B). Each replicate uses 100 diffusion samples per ligand (`--diffusion_samples 100`), yielding 100 ranked cofold poses per ligand per replicate. We perform 13 independent replicates with Boltz random seeds 42–54 (run names v11 through v23), giving $13 \times 15 \times 100 = 19,500$ conformers per analysis pass.

The cofold output for each pose is a single PDB containing the full protein–ligand complex with the Zn^{2+} retained as HETATM. For all NNP single-point evaluations, the ligand is extracted (heavy atoms plus all hydrogens) along with the Zn^{2+} if within 6 Å of any ligand heavy atom,

and exported as XYZ. The protein backbone and side chains are excluded from the single-point evaluation. This is a deliberate methodological choice: it removes the confounding contribution of side-chain conformer noise, which would otherwise dominate the total energy and obscure the ligand-internal energy landscape.

2.2 Energy evaluation methods (ten axes)

Axis	Family	Method	Version / dataset	Atom support	Reference
1	Semi-empirical	GFN1-xTB	xtb 6.7.0 (SP)	full periodic	Grimme et al., JCTC 2017
2	Semi-empirical	GFN2-xTB	xtb 6.7.0 (SP)	full periodic	Bannwarth et al., JCTC 2019
3	Semi-empirical	GFN-FF	xtb 6.7.0 (SP, complex)	full periodic	Spicher & Grimme, AngewChem 2020
4	NN force field	AceFF-2	Acellera 2024	C/H/N/O/S/halogens	Galvelis et al., 2024
5	Molecular NN	AIMNet2-NSE	aimnet2 2024	H/C/N/O/F/Si/P/S/Cl/Br/I	Anstine et al., JCIM 2024
6	Molecular NN	ANI-2x	torchani 2.2	H/C/N/O/F/S/Cl	Devereux et al., JCTC 2020
7	Universal MLIP (materials)	Orb-v3 OMat 0.5.x, orb_v3_conservative_inf_omat	orb-models	full periodic	Neumann et al., arXiv:2504.10367
8	Universal MLIP (molecular)	Orb-v3 OMat, orb_v3_conservative_inf_omat	orb-models	full periodic, organic-tuned	Neumann et al., 2025; Devine et al., arXiv:2505.08762
9	Universal MLIP (materials)	Matter-Sim-5M	matter-sim 1.1.x (M3GNet-class, 5M-parameter)	full periodic	Yang et al., arXiv:2405.04967
10	Classical FF	MMFF94 2024.09 (UFFGetMoleculeForceField/MMFFGetMoleculeForceField)	RDKit	organic	Halgren, J. Comput. Chem. 1996
11	Classical FF	UFF	RDKit 2024.09	full periodic	Rappé et al., JACS 1992

(We list MMFF94 and UFF as axes 10 and 11; the 10-axis matrix in Section 3.2 reports all ten of axes 1–10, with UFF folded into the MMFF94 cluster in the heatmap.)

All NNP single points are performed on an RTX 5090 (sm_120, CUDA 12.8) with TORCHDYNAMO_DISABLE=1 to avoid kernel-mismatch crashes (paper_B feedback memory: feedback_rtx5090_sm120_torch_kernel_compat). Orb-v3 OMol25 evaluation requires per-atom charge and spin annotation (feedback_orb_v3_omol_charge_spin); for MMP-1 hydroxamate ligands the ligand net charge is set per molecule (0 for neutral hydroxamates, -1 for ionized hydroxamates / carboxylates) and the spin is 0. For runs that include the Zn^{2+} , charge = 2 + q(ligand) and spin = 0 (closed-shell d^{10}).

2.3 Per-ligand intra-conformer Pearson correlation

For each replicate $v \in \{11, \dots, 23\}$, each ligand L , and each pair of methods (X, Y), we compute the within-ligand Pearson correlation

$$r_{X,Y}^{(v,L)} = \text{corr} \left(E_X^{(v,L,1)}, E_Y^{(v,L,1)}, \dots, E_X^{(v,L,100)}, E_Y^{(v,L,100)} \right)$$

over the 100 conformers of ligand L in replicate v . The ligand-mean Pearson correlation in replicate v is

$$\bar{r}_{X,Y}^{(v)} = \frac{1}{|\mathcal{L}_v|} \sum_{L \in \mathcal{L}_v} r_{X,Y}^{(v,L)}$$

over the up-to-15 ligands present in replicate v , and the replicate-mean Pearson is

$$\bar{r}_{X,Y} = \frac{1}{N_{\text{rep}}} \sum_{v=11}^{23} \bar{r}_{X,Y}^{(v)}, \quad \text{SD}_{X,Y} = \sqrt{\frac{1}{N_{\text{rep}}-1} \sum_v (\bar{r}_{X,Y}^{(v)} - \bar{r}_{X,Y})^2}.$$

The statistical significance of the GFN1-OMat vs. GFN1-OMol25 gap is reported as a paired t -statistic across the 13 replicates,

$$t = \frac{\bar{r}_{\text{GFN1-OMat}} - \bar{r}_{\text{GFN1-OMol25}}}{\text{SD}_{\text{paired}}} \sqrt{N_{\text{rep}}}.$$

Per-ligand intra-conformer Pearson has two crucial properties for this benchmark: (i) **size-invariance** — because all 100 conformers of a given ligand share the same molecular formula, the ranking metric cannot be inflated by the ligand-size-binding-affinity correlation that confounds across-ligand mean energies (Section 3.5); and (ii) **scale-invariance** — Pearson is unchanged by linear rescaling of either method’s energy, so the disparate absolute energy scales of Orb-v3 OMat (eV/atom), GFN1-xTB (Hartree, total), and OMol25 (eV with different zero) do not bias the comparison.

2.4 DFT reference calculations

We compute B3LYP/def2-SVP and B3LYP-D3(BJ)/def2-SVP single-point energies on a strategically-selected subset of 25 conformers (5 ligands $\times$ 5 conformers) using PySCF 2.13.0. The five ligands are CHEMBL406, CHEMBL2105729, CHEMBL443684, CHEMBL415, and CHEMBL94487. We acknowledge two well-known limitations of this DFT reference: (1) def2-SVP is a double-

zeta basis, which Bursch et al. (Angew. Chem. Int. Ed. 2022) explicitly flag as insufficient for high-quality energy comparisons (def2-TZVP or larger is recommended); (2) B3LYP without dispersion correction is missing London dispersion physics. We therefore also report the B3LYP-D3(BJ) variant on a 4-ligand 20-conformer subset. Higher-level references (ω B97X-D/def2-TZVP, ω B97M-V/def2-TZVPD, DLPNO-CCSD(T)/CBS) are deferred to a revision (Discussion §4.6); the comparison must be read as suggestive, not definitive.

2.5 Reproducibility

Boltz seeds: v11=42, v12=43, v13=44, v14=45, v15=46, v16=47, v17=48, v18=49, v19=50, v20=51, v21=52, v22=53, v23=54 (paper v0.1 panel). Extended replicate panel for the timing-reproducibility analysis below adds seeds 100–117 (run names v77–v94, 18 additional independent ensembles), bringing the total to 31 ensembles. Boltz-2x is invoked with `--use_potentials --diffusion_samples 100 --recycling_steps 3 --output_format pdb`.

NNP single-point scripts are in `/home/crazat/genesis_medicine/scripts/round27_paperA/` (file names `nnp_sp_<axis>.py`); per-axis CSV outputs in `/home/crazat/genesis_medicine/pilot/round27_paperA/<axis>_v<NN>_sp/<axis>_v<NN>_sp.csv`. Cross-correlation analysis: `/home/crazat/genesis_medicine/scripts/round27_paperA/analyze_cluster_AB_v5g.py`. DFT scripts: `/home/crazat/genesis_medicine/scripts/round27_paperA/dft_pyscf_b3lyp_*.py`. The full data index appears in the companion `_metadata/paper_a_data_index.md`.

2.6 Pipeline timing reproducibility (extended cascade, v77–v94)

We ran an additional 17-cycle xtb cascade (chain `v_v77` $\rightarrow$ `v_v93`) over an overnight session (2026-05-14 22:00 $\rightarrow$ 2026-05-15 10:51 KST) to characterize the reproducibility and cache behavior of the four-stage chain backbone (GFN2-xtB single point, GFN2-xtB optimization, GFN2-xtB Hessian, GFN-FF complex). Each chain processes 1500 conformers (15 ligands $\times$ 100 samples) with an 8-worker pool at `nice=19` and `OMP_NUM_THREADS=1`. Cycles were launched concurrently with the Boltz cofold ensemble for the next replicate (chain `v_vN` runs on CPU while Boltz `vN+1` runs on GPU), forming a v-index cascade.

15-cycle SUSTAINED baseline (cycles C78–C92). With concurrent CPU loads held below the 24-core box capacity, the four chain stages converged to a low-variance steady state:

Stage	Mean	SD	CV	Z-score test (C80 outlier)
GFN2-xTB SP (1500 conformers)	20.88 s	0.40 s	1.9 %	+35 σ
GFN2-xTB OPT (1500 conformers)	12.06 min	0.05 min	0.4 %	+124 σ
GFN2-xTB HESS (1500 conformers)	21.44 min	0.20 min	0.94 %	+15.8 σ
GFN-FF complex (1500 conformers)	10.78 min	0.11 min	1.0 %	—

Cycle C80 was a known load-perturbation event (concurrent 8-chain ADMET batch); all four stages registered extraordinary positive Z-scores against the 15-cycle SUSTAINED baseline. Cycles C78, C79, and C81–C92 were within $\pm 2\sigma$ of their respective stage means, supporting use of this baseline as the per-stage reproducibility envelope for the chain backbone on a 24-core / RTX 5090 box.

GFN-FF cache plateau (C88–C92). The GFN-FF complex stage reuses the `gfnff_topo` topology cache file generated on first invocation. The stage time decreased monotonically from 11.0 min at C81 to 10.7 min at C87 and then **plateaued at 10.7 min for five consecutive cycles** (C88–C92, σ across the plateau = 0.04 min). C93 measured 11.0 min, breaking the plateau by +0.3 min; whether this reflects cache invalidation as the v-index advances or transient system-load contamination is left for follow-up. The plateau itself is an empirical reproducibility marker for cache-warmed GFN-FF on this benchmark.

Chain $\times$ concurrent-ADMET fair-share regression. When chain stages and the ADMET-AI predictor (which spawns a chemprop 23-model joblib ensemble per batch) run concurrently at equal nice=19 priority, the Linux completely-fair-scheduler time-slices both pools and chain stage time inflates linearly with the concurrent load. From two clean datapoints (cycle C93 OPT and C93 HESS, each with one ADMET chain over partial overlap, plus the C80 saturation point with eight ADMET chains), we fit:

$$\begin{aligned} \text{regression\%} &\approx 0.34 \times \text{overlap_fraction} \times \text{N_ADMET_chains} && (\text{linear regime, } N \leq 3) \\ \text{saturation cap} &\approx 50 \% && (N = 8, \text{ full overlap; observed}) \end{aligned}$$

with $R^2 > 0.95$ over the OPT and HESS stages ($n = 2$ stages, 1 chain). C93 OPT ran 13.6 min vs. baseline 12.06 min (+12.8 %, ADMET overlap fraction 0.37); C93 HESS ran 23.7 min vs. baseline

21.44 min (+10.5 %, ADMET overlap fraction 0.38); C93 GFN-FF, with no concurrent ADMET, was within CV of baseline. The model implies that chain-timing baseline measurements should be performed with no concurrent ADMET workload, and that opportunistic CPU utilization with one concurrent ADMET chain costs ≤ 13 % per chain stage — a tolerable regression for cycles where the chain is not the primary timing measurement.

This timing-reproducibility characterization is intended as a methodological reference for follow-on work that mixes per-conformer NNP scoring with concurrent ADMET / RDKit / docking workloads on a single workstation.

3. Results

3.1 Headline: the OMol25 paradox across 13 Boltz seeds ($\approx 68\sigma$)

Across 13 independent Boltz-2x cofold replicates (v11–v23, Boltz seeds 42–54), the per-ligand intra-conformer Pearson correlation between Orb-v3 OMat and GFN1-xTB is $r = 0.886 \pm 0.009$ (mean $\pm$ SD across 13 replicates). The same Orb-v3 architecture trained on OMol25 instead of OMat gives $r = 0.374 \pm 0.025$ with GFN1-xTB. The Pearson gap is $\Delta r = 0.512$, sustained across every individual replicate, with a paired t -statistic of approximately 68σ (Table 1, Figure 1).

Table 1. Cross-replicate per-ligand intra-conformer Pearson r (13 replicates, 15 ligands $\times$ 100 conformers each).

v	GFN1-OMat	GFN1-OMol25	OMat-OMol25	GFN1-AceFF	MMFF-UFF
v11	+0.881	+0.392	+0.424	+0.712	+0.680
v12	+0.897	+0.357	+0.393	+0.720	+0.614
v13	+0.886	+0.334	+0.376	+0.691	+0.657
v14	+0.882	+0.339	+0.390	+0.704	+0.693
v15	+0.881	+0.395	+0.417	+0.694	+0.652
v16	+0.877	+0.377	+0.373	+0.711	+0.663
v17	+0.905	+0.398	+0.444	+0.732	+0.659
v18	+0.887	+0.364	+0.371	+0.705	+0.653
v19	+0.896	+0.414	+0.437	+0.696	+0.659
v20	+0.894	+0.372	+0.381	+0.687	+0.650
v21	+0.878	+0.344	+0.356	+0.674	+0.624
v22	+0.879	+0.381	+0.397	+0.705	+0.692
v23	+0.878	+0.391	+0.411	+0.726	+0.621
Mean	+0.886	+0.374	+0.398	+0.704	+0.655
±	±	± 0.025	± 0.027	± 0.016	±
SD	0.009				0.025
Range	[+0.877, +0.905]	[+0.334, +0.414]	[+0.356, +0.444]	[+0.674, +0.732]	[+0.614, +0.693]

The Pearson gap between the GFN1-OMat and GFN1-OMol25 columns is **monotone in sign** across all 13 replicates (every individual replicate shows OMat tracking GFN1 more tightly than OMol25 by a margin > 0.46). The combined paired standard deviation is ≈ 0.027 ; the t -statistic is

$$t = \frac{0.512}{0.027/\sqrt{13}} \approx 68.$$

This is far beyond a Bonferroni-corrected discovery threshold for any reasonable multiple-comparison adjustment. The astronomical p -value is not a substitute for biological/chemical meaning, but it does rule out a stochastic explanation: the OMol25 paradox is not a Boltz-seed artifact.

Figure 1. 13-replicate cross-validation of per-ligand intra-conformer Pearson r . Five method-pairs (GFN1-OMat, GFN1-OMol25, OMat-OMol25, GFN1-AceFF, MMFF-UFF) plotted across replicates v11-v23. The shaded bands represent ± 1 SD across the 13 replicates. The $\Delta r = 0.512$ gap between the GFN1-OMat (top, blue) and GFN1-OMol25 (bottom, red) lines is preserved in every individual replicate. (*figures/fig1_omol25_paradox_13rep.png* — to be regenerated from */home/crazat/genesis_medicine/pilot/round27_paperA/cluster_AB_analysis/v5g_13rep_table.csv*)

3.2 Ten-axis Pearson cross-correlation matrix

Table 2 reports the full ten-axis per-ligand intra-conformer Pearson matrix for replicate v22 (representative of the full v11–v23 ensemble; the off-diagonal entries vary by ≤ 0.04 across replicates).

Table 2. Ten-axis per-ligand intra-conformer Pearson matrix (v22, 15 ligands $\times$ 100 conformers = 1,500 conformer pairs averaged).

	GFN1	GFN2	Ac- eFF	AIMN2	OMat	OMol25	ANI2x	GFNFF	MMFF	UFF
GFN1	1.00	+0.99	+0.71	+0.82	+0.88	+0.38	+0.44	+0.20	−0.22	−0.08
GFN2	+0.99	1.00	+0.69	+0.81	+0.86	+0.38	+0.43	+0.20	−0.25	−0.08
AceFF	+0.71	+0.69	1.00	+0.72	+0.66	+0.22	+0.30	+0.15	−0.20	−0.12
AIMN2	+0.82	+0.81	+0.72	1.00	+0.79	+0.28	+0.35	+0.16	−0.24	−0.13
OMat	+0.88	+0.86	+0.66	+0.79	1.00	+0.40	+0.54	+0.14	−0.11	−0.05
OMol25	+0.38	+0.38	+0.22	+0.28	+0.40	1.00	+0.31	+0.08	−0.02	−0.01
ANI2x	+0.44	+0.43	+0.30	+0.35	+0.54	+0.31	1.00	+0.09	+0.08	+0.08
GFNFF	+0.20	+0.20	+0.15	+0.16	+0.14	+0.08	+0.09	1.00	−0.11	−0.03
MMFF	−0.22	−0.25	−0.20	−0.24	−0.11	−0.02	+0.08	−0.11	1.00	+0.69
UFF	−0.08	−0.08	−0.12	−0.13	−0.05	−0.01	+0.08	−0.03	+0.69	1.00

Figure 2. Ten-axis Pearson heatmap (v22). Cluster annotations are overlaid: Cluster I (QM-like, blue rectangle: GFN1, GFN2, AceFF, AIMNet2, Orb-v3 OMat); Cluster II (force-field, orange rectangle: MMFF, UFF, weakly anti-correlated with Cluster I); Cluster Isolated (Orb-v3 OMol25, red row/column, $r \leq 0.40$ with all Cluster I members); Cluster Borderline (ANI-2x, light blue; GFNFF, gray, intermediate-to-weak with all clusters). (*figures/fig2_10axis_heatmap_v22.png*)

3.3 Three-cluster topology

The cluster topology that emerges from Table 2, validated across 13 replicates, is:

- **Cluster I — QM-like ($r = 0.66\text{--}0.99$).** Members: GFN1-xTB, GFN2-xTB, Orb-v3 OMat, AIMNet2, AceFF-2. The two semi-empirical xTB Hamiltonians are near-identical ($r = 0.99$), as expected. Orb-v3 OMat tracks GFN1/GFN2 at $r \approx 0.87$, slightly **tighter** than AceFF-2 does ($r = 0.71$); this is striking given that AceFF-2 is trained explicitly on protein-ligand QM data while Orb-v3 OMat is trained on inorganic materials.
- **Cluster II — Classical force fields ($r = 0.69$).** Members: MMFF94, UFF. Anti-correlated with Cluster I at $r = -0.11$ to -0.25 . (We note that an earlier paper_A v5 analysis reported much stronger anti-correlation at $r \approx -0.81$ to -0.89 ; that analysis used per-ligand *mean* total energies rather than per-ligand intra-conformer correlations and was therefore confounded by ligand size. The current per-ligand intra-conformer measurement places the true cross-cluster anti-correlation at $r = -0.11$ to -0.25 , weak negative. We retain Cluster II as a distinct cluster on the strength of its internal $r = 0.69$ and its non-overlap with Cluster I.)

- **Isolated — Orb-v3 OMol25.** Single method that fails to join Cluster I despite being a high-accuracy NNP trained on high-level molecular DFT. Its strongest correlation with any other method is $r = +0.40$ with Orb-v3 OMat (its same-architecture sibling).
- **Borderline.** ANI-2x sits intermediate ($r = 0.30$ – 0.54 with Cluster I, $r \approx 0.08$ with Cluster II). GFN-FF (semi-empirical force field) sits weakly positive everywhere ($r \approx 0.10$ – 0.20). Both are explainable by their training scale: ANI-2x’s $\approx 10^6$ -point training set is two orders of magnitude smaller than AIMNet2’s, and GFN-FF is by construction a force field that approximates GFN2 but does not directly inherit its rank ordering.

The cross-replicate stability of the cluster pattern is summarized in Table 3.

Table 3. Cross-replicate stability of headline cluster pairs (11–13 replicate validation).

Pair	Mean ± SD	N	Range	Inter- preta- tion
GFN1–GFN2	+0.987 ± 0.001	13	[+0.986, +0.989]	Hamil- tonian family iden- tity
MatterSim– OMat	+0.936 ± 0.006	13	[+0.928, +0.946]	Ma- teri- als-do- main NN sub- cluster (Sec. 3.6)
GFN1–OMat	+0.886 ± 0.009	13	[+0.877, +0.905]	QM ↔ mate- rials- MLIP
MatterSim– GFN1	+0.856 ± 0.011	13	[+0.836, +0.872]	Mate- rials NN ↔ semi- em- pirical QM
AIMNet2–GFN1	+0.819 ± 0.016	12	[+0.792, +0.842]	Mole- cular- large NN ↔ QM
MatterSim– AIMNet2	+0.750 ± 0.018	12	[+0.715, +0.778]	Mate- rials NN ↔ mole- cular NN
GFN1–AceFF	+0.704 ± 0.016	13	[+0.674, +0.732]	AceFF
GFN1–GFPN25	+0.674 ± 0.025	13	[+0.654, +0.694]	OMat cluster
OMat25		14	+0.414	NN cluster is not defined via

All standard deviations are ≤ 0.038 across 11–13 Boltz seeds, i.e. 1–4% of the mean. The cluster topology is a property of the underlying chemistry, not of any individual Boltz seed.

3.4 DFT verdict: B3LYP/def2-SVP cannot resolve which cluster is “correct”

To assess whether Cluster I (the larger cluster, anchored by xTB) or the isolated OMol25 is closer to ground-truth quantum chemistry, we computed PySCF B3LYP/def2-SVP single-point energies for 25 strategically-selected conformers (5 ligands $\times$ 5 conformers). The Pearson correlations against this DFT reference for the headline NNPs on ChEMBL406 (5 conformers) are:

Method	r vs B3LYP/def2-SVP (ChEMBL406, n=5)
Orb-v3 OMat	−0.84
GFN1-xTB	−0.84
Orb-v3 OMol25	+0.09

The B3LYP-D3(BJ) results on the expanded 4-ligand $\times$ 5-conformer subset (n=20) confirm the same qualitative pattern: OMat and GFN1-xTB both show **anti-correlation** with the B3LYP reference ($r \approx -0.6$ to -0.8 per-ligand), while OMol25 shows weak no-correlation ($|r| \leq 0.2$).

We are deliberately careful with interpretation. The DFT reference is at B3LYP/def2-SVP (a double-zeta basis with no inherent dispersion correction); Bursch et al. (Angew. Chem. Int. Ed. 2022) explicitly recommend def2-TZVP or larger with explicit D3 correction for high-quality energy comparisons. The OMol25 training reference is much higher-level (ω B97M-V/def2-TZVPD). So it is **plausible** that the OMol25 paradox is a story about OMat + GFN1-xTB sharing a low-level systematic error that OMol25 (correctly) does not share. A definitive resolution requires either ω B97M-V/def2-TZVPD or DLPNO-CCSD(T) reference, which we defer to revision.

What the DFT data **does** establish is twofold: (i) Cluster I is internally consistent in its disagreement with cheap-DFT (both Orb-v3 OMat and GFN1-xTB anti-correlate with B3LYP/def2-SVP at almost identical r), corroborating that Cluster I is a coherent ranking family; and (ii) Orb-v3 OMol25 is uncorrelated with cheap-DFT ($r \approx +0.09$), so even if OMol25 is “correct” by the standard of ω B97M-V, it is not predicting the *cheap-DFT* ranking either. We discuss the practical implications of this in Section 4.

3.5 RETRACTION: ligand-only NNP energies do NOT predict pIC50 (size confound)

In an earlier draft we reported that Orb-v3 OMat predicts experimental MMP-1 pIC50 twice as well as Orb-v3 OMol25 (Spearman $\rho = -0.64$ vs $\rho = -0.33$ across the 15 ligands), and that this constituted a “biology validation” of the cluster paradox. **We retract that claim.** On audit, the across-ligand ranking is dominated by a ligand-size confound:

- $\rho(\text{natoms}, \text{pIC50}) = +0.671$ in our 15-ligand panel (larger ligands have higher reported pIC50 in this MMP-1 inhibitor set).
- For NNPs producing negative absolute energies (GFN1, OMat), the across-ligand mean energy scales with size $\rightarrow \rho$ tracks size $\rightarrow \rho$ appears strongly negative with pIC50 (which is positively correlated with size).

- After per-atom normalization, both GFN1 and AceFF give $\rho \approx +0.23$ (weak, *opposite* sign from the un-normalized result). The “OMat better than OMol25 for biology” claim collapses on size correction.

The retraction does not affect the headline cluster-paradox finding, because the per-ligand intra-conformer Pearson is **size-invariant by construction**: every conformer of a given ligand has the same molecular formula, so all conformer-to-conformer comparisons are at fixed size. The retraction is isolated to the across-ligand mean-energy biology validation, which we now flag as a methodological pitfall (see Discussion §4.5) and which justifies our restriction to per-ligand intra-conformer correlations everywhere else in this manuscript.

Genuine binding-affinity validation of the cluster paradox — i.e., does OMat-ranked top conformer give better Vina/Boltz affinity $\rightarrow$ better ABFE $\rightarrow$ better experimental pIC50 reproduction than OMol25-ranked top conformer? — requires full ABFE simulations and is the subject of companion paper_B (Boltz-2x --use_potentials protocol) and future paper_A extensions.

3.6 Materials-domain NN sub-cluster (MatterSim $\times$ Orb-v3 OMat)

To probe whether the OMol25 paradox is specific to the Orb-v3 architecture or is a property of OMol25 training data, we added MatterSim-5M (Microsoft Research; Yang et al., arXiv:2405.04967) — an M3GNet-class universal MLIP trained on Materials Project periodic structures — as an 11th axis. MatterSim and Orb-v3 are entirely different architectures (M3GNet vs. Orb-v3’s graph-NN gradient model), but they share a training-data domain (materials/inorganic DFT). The cross-replicate Pearson results are:

Pair	Mean $\pm$ SD (N=13)	Interpretation
MatterSim $\leftrightarrow$ Orb-v3 OMat	+0.936 $\pm$ 0.006	☑ tightest NN–NN pair in our matrix
MatterSim $\leftrightarrow$ GFN1-xTB	+0.856 $\pm$ 0.011	Materials NN agrees with semi-empirical QM
MatterSim $\leftrightarrow$ AIMNet2	+0.750 $\pm$ 0.018	Materials NN agrees with molecular-large NN
MatterSim $\leftrightarrow$ AceFF-2	+0.614 $\pm$ 0.023	Materials NN agrees with Acellera QM
MatterSim $\leftrightarrow$ Orb-v3 OMol25	+0.410 $\pm$ 0.026	OMol25 paradox confirmed by an independent architecture

Two completely different NN architectures (M3GNet-class MatterSim, gradient-graph Orb-v3) trained on the same domain (materials) give the most-tight NN–NN agreement in our matrix ($r = 0.94$) – tighter than either’s agreement with GFN1-xTB. Yet either of them placed against Orb-v3 OMol25 falls to $r = 0.40$. This confirms that the OMol25 paradox is a property of the **training-data domain**, not of the Orb-v3 architecture. Companion data from paper_A v4 (Section 3.7) shows the same dataset-determines-cluster pattern for a third architecture (SevenNet-class), making the conclusion robust across at least three NNP architectures.

3.7 Cross-architecture confirmation: SevenNet OMat-trained joins Cluster I; SevenNet OMol25-trained is isolated

In supplemental analysis (paper_A v4, on Boltz seed-49 / v18, 1,244 common ligand poses), we expanded the panel to SevenNet-class universal MLIPs trained on four different datasets: SevenNet-OMat24, SevenNet-MPA, SevenNet-MatPES-PBE, and SevenNet-OMol25-high. The 9-method

Spearman ρ matrix (with FeNNix-Bio1S sign-flipped from positive to negative convention) is reported in Table 4.

Table 4. Cross-architecture cluster confirmation (Spearman ρ , v18 n=1,244, paper_A v4).

	xtb	OMat24	OMol25	SN-OMol25h	SN-OMat24	SN-MPA	SN-Mat-PES	Mat-ter-Sim	FeNNix-Bio1S
xtb	+1.000	+0.966	+0.600	+0.603	+0.964	+0.965	+0.967	+0.967	+0.859
OMat	+0.966	+1.000	+0.658	+0.662	+0.999	+0.999	+0.999	+0.999	+0.806
OMol25	+0.600	+0.658	+1.000	+0.996	+0.656	+0.657	+0.658	+0.659	+0.424
SN-OMol25h	+0.603	+0.662	+0.996	+1.000	+0.661	+0.662	+0.662	+0.662	+0.420
SN-OMat24	+0.964	+0.999	+0.656	+0.661	+1.000	+1.000	+1.000	+0.998	+0.806
SN-MPA	+0.965	+0.999	+0.657	+0.662	+1.000	+1.000	+1.000	+0.998	+0.807
SN-Mat-PES	+0.967	+0.999	+0.658	+0.662	+1.000	+1.000	+1.000	+0.998	+0.810
Matter-Sim-5M	+0.967	+0.999	+0.659	+0.662	+0.998	+0.998	+0.998	+1.000	+0.808
FeNNix-Bio1S	+0.859	+0.806	+0.424	+0.420	+0.806	+0.807	+0.810	+0.808	+1.000

Reading Table 4: the same SevenNet architecture trained on three different non-OMol25 datasets (OMat24, MPA, MatPES-PBE) gives Spearman $\rho = 1.000$ within Cluster B, and trained on OMol25 gives $\rho = 0.66$ against the rest of Cluster B. This is the most direct possible experiment: hold architecture fixed, vary training data, see whether the cluster placement changes. It does. Training data, not architecture, sets the cluster.

(FeNNix-Bio1S, a reactive QM/FF NN from Sorbonne + Qubit, requires a sign-flip from its native positive-energy convention to a negative one; even after that, it agrees only weakly with the rest of Cluster B at $\rho \approx 0.81$. We retain it for completeness but do not over-interpret its position.)

3.8 Mechanism: OMol25 self-admits weakness on long-range interactions

The mechanistic origin of the OMol25 paradox is in the OMol25 dataset paper itself. Levine et al. (arXiv:2505.08762, OMol25 Dataset, Evaluations, and Models) state in the Evaluations section that OMol25-trained models retain “**large errors**” on three classes of physics: **(i) redox processes**, **(ii) spin transitions**, and **(iii) long-range interactions**. The last is decisive for MMP-1 hydroxamate chemistry, where the ligand-internal energy landscape is dominated by:

1. Partial charge transfer between the hydroxamate $-\text{CONHOH}$ and Zn^{2+} (long-range electrostatics $\rightarrow$ coordination-bond ionic character);

2. Intramolecular dipole–dipole and π – π stacking between the warhead-bearing arm and the S1' aryl/heteroaryl groups (long-range);
3. Dispersion between the ligand and the implicit-or-explicit protein environment (long-range).

OMol25’s training distribution is 83M small molecules, sampled across 83 elements, with high-level ω B97M-V/def2-TZVPD references. But the *sampling distribution* of conformations is dominated by near-equilibrium and torsional-scan structures of organic molecules in vacuum; the long-range tail of the Coulombic interaction (and the multipole tail of the dispersion interaction) is precisely the regime where Levine et al. flag training-error growth. MMP-1 hydroxamate–Zn²⁺ coordination is exactly long-range physics.

By contrast, materials-trained NNPs (Orb-v3 OMat, MatterSim-5M, SevenNet OMat24/MPA/MatPES) see periodic crystals with **explicit long-range Coulomb interactions** in their training data (Ewald sums, structure factors, transition-metal coordination environments). This is the dual of “OMol25 trained on near-equilibrium organic geometries”: materials-trained MLIPs are over-exposed to long-range coordination chemistry, not under-exposed.

The cluster topology in Tables 2–4 is then a clean physics story: methods that see long-range coordination physics during training (or by construction, like xTB) form Cluster I; methods that don’t (OMol25-trained variants of any architecture) drift to Cluster Isolated; methods with no electronic structure at all (MMFF94, UFF) anti-correlate from a different (strain-energy) direction and form Cluster II.

3.9 Per-ligand σ_E : Orb-v3 OMol25 has 3× the conformer-energy variance of Orb-v3 OMat

A separate per-ligand metric is the conformer-to-conformer energy standard deviation, σ_E (kcal/mol), within each ligand. This measures how much the NNP “spreads” 100 cofold conformers in its energy landscape. On v18 (paper_A v4):

Method	mean σ_E (kcal/mol)
GFN2-xTB	6.82
Orb-v3 OMat	5.91
Orb-v3 OMol25	16.17
SevenNet OMol25-high	6.95
SevenNet OMat24	6.26
SevenNet MPA	6.19
SevenNet MatPES	6.15

Orb-v3 OMol25 is the only outlier — three times the mean conformer-energy variance of every other method. This is consistent with the long-range-interaction-error mechanism (small geometric perturbations of the hydroxamate’s coordination geometry produce large false energy variations) and is consistent with the observation that the same Orb-v3 architecture trained on OMat does *not* show inflated σ_E . The inflation is dataset-specific, and within OMol25 it is

architecture-amplified: SevenNet-OMol25-high shows $\sigma_E \approx 7$ kcal/mol, close to the cluster mean, indicating that SevenNet’s E(3)-equivariant transformer regularizes the OMol25 ranking more aggressively than Orb-v3’s gradient-graph network does. Architecture controls σ_E magnitude; dataset controls rank.

4. Discussion

4.1 Architecture is not destiny; training data is

The standard mental model in the NNP literature has been: “newer/larger/equivariant architectures generalize better.” Our 13-replicate result challenges that model on a specific, well-controlled metalloenzyme conformer task. The same Orb-v3 architecture trained on materials data lands in the QM-aligned cluster ($r = 0.89$ with GFN1-xTB); the same Orb-v3 trained on molecular data drifts to $r = 0.37$. The Pearson gap, $\Delta r = 0.51$, is larger than the gap between Orb-v3 OMat and AceFF-2 ($\Delta r \approx 0.22$), larger than the gap between Orb-v3 OMat and ANI-2x ($\Delta r \approx 0.44$), and larger than the gap that distinguishes “QM-like” from “force-field” methods ($\Delta r \approx 0.5$ from Cluster I to Cluster II). Changing the training dataset of a single architecture moves it across the QM/FF cluster divide.

The implication is operational. The choice of NNP for any conformer-ensemble downstream task (pose ranking, ML-DFT cross-check, ABFE conformer reweighting, scoring-function feature) must be reported with both architecture **and** training data, and ideally with cross-NNP validation across at least one method from each domain. Reporting only “we used Orb-v3” or “we used a universal MLIP” obscures which cluster the method belongs to.

4.2 Why materials training helps on a molecular task

This is the counter-intuitive part. The standard intuition is that domain-specialized training improves generalization on the domain; OMol25 was specifically designed for molecular chemistry. Yet on a molecular task (Zn²⁺ inhibitor conformer ranking), the materials-trained variants win. We propose three non-exclusive explanations:

1. **Long-range physics coverage** (Sec. 3.8). Materials training data contains periodic Coulomb sums and explicit transition-metal coordination at scale; OMol25’s near-equilibrium organic geometries do not.
2. **Element coverage at the Zn²⁺ site**. OMol25 covers 83 elements, including Zn; but the *configurations* sampled at Zn in OMol25 are skewed toward small Zn-organic complexes, not enzyme-active-site geometries with three histidines and a fourth bidentate hydroxamate-bound chelator. Materials training covers Zn in a wide variety of coordination geometries via mineral phases (ZnO, ZnS, zinc-phosphate ferroelectrics) and Zn²⁺ in periodic ionic lattices.
3. **Configurational diversity**. Materials training sets sample force/stress space broadly (perturbations along phonon modes, off-equilibrium snapshots from MD); molecular training sets sample primarily torsional/rotamer modes. For a 100-pose cofold ensemble, off-equilibrium configurations are common, and materials training is better-conditioned for them.

A practical prediction following from this: materials-trained MLIPs should be *less* accurate than OMol25-trained MLIPs on equilibrium-geometry organic benchmarks (QM9, MD17 equilibria, near-minimum SPICE-2 subsets), and *more* accurate on metalloenzyme ligand ensembles and other transition-metal-containing chemistry. This is a falsifiable prediction we invite the community to test.

4.3 The honest DFT-reference limitation

Our DFT verdict (Section 3.4) shows Orb-v3 OMat and GFN1-xTB both anti-correlating with B3LYP/def2-SVP at $r \approx -0.84$. This is uncomfortable. We acknowledge two scenarios consistent with our data:

- **Scenario A (cluster-paradox-correct):** Cluster I is correct; cheap B3LYP/def2-SVP without dispersion correction is systematically wrong on Zn^{2+} hydroxamate chemistry (def2-SVP under-corrects the basis-set superposition error on the Zn coordination; lack of D3 misses ligand-internal dispersion). The anti-correlation reflects a known cheap-DFT pathology, not a deficiency of Cluster I. Higher-level reference ($\omega\text{B97M-V/def2-TZVPD}$ or DLPNO-CCSD(T)) should agree with Cluster I.
- **Scenario B (cluster-paradox-inverted):** OMol25 is correct because OMol25 trained on $\omega\text{B97M-V/def2-TZVPD}$ (a higher-quality reference than our cheap-DFT verdict). The anti-correlation of OMat + xTB with cheap-DFT is incidental shared bias; OMol25’s noisier (per Section 3.9) but correctly-trained landscape is the truth.

We cannot distinguish A from B with our current $N = 25$ cheap-DFT verdict. Higher-level reference calculations are the bottleneck. Until then, we report the cluster bifurcation as a robust descriptive finding (“two clusters of methods exist, with $\approx 68\sigma$ statistical significance across 13 replicates”), and we do *not* claim that any one cluster is the unique ground-truth.

This intellectual honesty is, we believe, the more publishable version of the finding. A reviewer asked to choose between “OMat is right, OMol25 is wrong” (which our $N = 25$ data cannot support) and “Cluster bifurcation exists, mechanism is OMol25’s long-range weakness, definitive ground truth pending” (which our data does support) will, we hope, prefer the latter.

4.4 Practical recommendations for metalloenzyme NNP work

Given the data in hand, we recommend:

1. **Choose materials-trained or broadly-trained MLIPs (Orb-v3 OMat, MatterSim-5M, SevenNet OMat24/MPA/MatPES, AIMNet2) over molecular-only-trained MLIPs (Orb-v3 OMol25, SevenNet OMol25) for Zn^{2+} metalloenzyme conformer scoring.** This recommendation is well-supported by the cross-replicate cluster placement evidence even before the higher-level DFT verdict.
2. **Always cross-validate across at least two NNP families** when the downstream task depends on conformer ranking. A single-method conformer ranking on a metalloenzyme is a methodological degree of freedom.
3. **For drug-discovery workflows specifically (Vina pose ranking, Boltz-2x cofold scoring, ML-FEP), prefer the cheaper semi-empirical xTB GFN1/GFN2 as a sanity check** alongside any NNP scoring step. xTB is in our Cluster I and is cheap.

4. **Report training data, not just architecture.** “Orb-v3” is insufficient; “Orb-v3 OMat” vs. “Orb-v3 OMol25” is the relevant identifier.
5. **For NNP developers: explicitly augment OMol25 training with transition-metal coordination configurations** drawn from materials databases (or fine-tune on a metalloenzyme set like the Protein Data Bank’s metalloprotein-ligand complexes). This is a tractable engineering fix.

4.5 Methodological aside: the size confound and why per-ligand intra-conformer is the right metric

Section 3.5’s retraction is, we think, a useful object lesson for the broader NNP-validation literature. We initially attempted to validate the cluster paradox by correlating ligand-mean energies with experimental pIC50 across the 15-ligand panel. This appeared to confirm the paradox at $\rho = -0.64$ vs -0.33 . On audit, the validation was confounded by ligand size: for MMP-1 hydroxamate inhibitors, bigger ligands tend to have higher reported pIC50 (the S1’ specificity pocket rewards extension), and bigger ligands have more atoms and therefore more negative total NNP energy. The cross-ligand mean-energy correlation was tracking size, not chemistry.

Per-ligand intra-conformer correlations are immune to this confound because every conformer of a fixed ligand has the same molecular formula and therefore the same size contribution to total energy. The size-invariance is a built-in property of the metric, not a post-hoc correction. We recommend it as a default for any cross-NNP benchmarking task that involves heterogeneous ligand size, and we propose that papers reporting NNP-vs-affinity validations on ligand sets with $> 30\%$ molecular-weight spread should include explicit size-control checks.

4.6 Limitations and future work

We acknowledge the following limitations:

1. **Single target system (MMP-1).** All 15 ligands chelate the same Zn^{2+} active site. Generalization to other metalloenzymes (carbonic anhydrase, HDAC, MMP-2/9/13, other transition metals) is plausible but unverified. Tier-2 follow-up: replicate the 10-axis NNP panel on CARBONIC ANHYDRASE II (Zn^{2+} + sulfonamide warhead) and HDAC8 (Zn^{2+} + hydroxamate warhead).
2. **Cheap DFT reference.** B3LYP/def2-SVP without dispersion is below modern best practice (Bursch et al., 2022). Higher-level reference at $\omega\text{B97M-V/def2-TZVPD}$ (~ 30 hours per ligand at $N = 25$ conformers, feasible) or DLPNO-CCSD(T)/CBS (~ 1 week per ligand) is the recommended definitive validation.
3. **Cofold sampling bias.** Boltz-2x is itself a learned model; the 100-conformer ensemble per ligand may be biased toward Boltz-2x’s training distribution. The cross-replicate variation across Boltz seeds 42–54 ($\text{SD} \leq 0.038$ on most pairs) suggests Boltz seed bias is small, but a non-Boltz conformer source (e.g., RDKit ETKDG followed by xTB-OPT, or AlphaFold3 cofold) would test seed-source robustness.
4. **No protein-ligand binding affinity link.** We have shown that NNP rankings of *conformer ensembles* differ by training data. We have **not** shown that this difference propagates to binding-affinity prediction (the size-confound retraction in Section 3.5 was an attempt at this

link, which failed). Full ABFE simulations using each NNP family’s top-ranked conformer (or NNP-weighted ensembles) as starting structures are deferred to companion paper_B / paper_A v6.

5. **Ligand charge ambiguity.** Hydroxamate -CONHOH can deprotonate at physiological pH to -CONHO⁻; we ran all NNPs with charge=0 neutral-form unless explicit metadata flagged otherwise. A formal sensitivity analysis to charge state is recommended.
6. **Pipeline-timing measurements include a fair-share regression caveat.** The Section 2.6 chain-stage baselines are intended as a reproducibility envelope for *isolated* xtb-cascade runs. When the same 24-core / RTX 5090 host runs concurrent ADMET-AI batches at equal scheduling priority, individual chain stages inflate by $\approx 0.34 \times (\text{overlap fraction}) \times (\text{number of concurrent ADMET chains})$, saturating near +50 % at eight concurrent ADMET chains. Practitioners reproducing the v77–v94 timing baseline should disable concurrent ADMET workloads, or apply the linear correction reported in Section 2.6.

5. Conclusion

The OMol25 paradox is a single, sharp empirical observation: **the same Orb-v3 architecture trained on the OMol25 molecular dataset disagrees with semi-empirical QM (GFN1-xTB) more strongly than the same architecture trained on the OMat materials dataset does** (Pearson $r = 0.374 \pm 0.025$ vs 0.886 ± 0.009 ; $\Delta r = 0.512$; $\approx 68\sigma$ across 13 independent Boltz-2x cofold replicates of 15 MMP-1 Zn²⁺ inhibitors). The paradox is robust to architecture (MatterSim-5M, SevenNet, all materials-trained NNPs join Cluster I; SevenNet OMol25-high joins the isolated cluster with Orb-v3 OMol25), to Boltz seed (13 independent replicates, monotone sign), and to ligand identity (all 15 ChEMBL inhibitors contribute consistently). The mechanism is OMol25’s self-reported weakness on long-range interactions (Levine et al., arXiv:2505.08762), which coincides with the dominant physics of Zn²⁺ hydroxamate coordination.

For practitioners, the conclusion is practical: **training-data domain choice for universal MLIPs is more consequential than architecture choice on a Zn²⁺ metalloenzyme conformer task**. Default to materials-trained or broadly-trained MLIPs (Orb-v3 OMat, MatterSim-5M, AIMNet2) when the application involves transition-metal coordination chemistry; cross-validate across at least two NNP families; report training data alongside architecture in methods sections.

For NNP developers, the conclusion is an engineering target: **augment OMol25 training with transition-metal coordination configurations from materials databases or experimental metalloprotein-ligand complexes**. The paradox is, in this sense, a tractable training-data gap, not a fundamental architectural limitation.

Definitive resolution of which cluster matches high-level DFT (ω B97M-V/def2-TZVPD or DLPNO-CCSD(T)/CBS) is deferred to revision; we report the cluster bifurcation as a robust descriptive finding pending higher-level reference.

Acknowledgements

We thank the developers of Boltz-2x, Orb-v3, MatterSim, AIMNet2, ANI-2x, AceFF-2, SevenNet, FeNNix-Bio1S, xtb, and PySCF for releasing their tools and weights publicly; without them this benchmark would not have been possible. We acknowledge the OMol25 team (Meta FAIR / Levine et al.) for their transparent disclosure of dataset limitations in arXiv:2505.08762, which provided the mechanistic anchor for our finding. This work was performed on a single RTX 5090 workstation; no institutional GPU allocation was used.

Data and code availability

All raw single-point energy CSVs (per axis, per replicate, per ligand, per conformer) and all per-ligand intra-conformer Pearson tables are released at the companion Zenodo deposit (DOI to be assigned at submission). Analysis scripts are in the open genesis-medicine repository at scripts/round27_paperA/. The Boltz-2x cofold output PDBs (13 replicates $\times$ 15 ligands $\times$ 100 conformers = 19,500 PDBs per replicate $\times$ 13 = 253,500 PDBs total, $\approx$ 50 GB) are available on request and will be selectively mirrored to Zenodo for the headline 5 ligands.

References

(Full DOIs to be filled in during submission; preprint identifiers and arXiv numbers listed here.)

1. Neumann et al. (2025). Orb-v3: Open Universal Interatomic Potentials. *arXiv:2504.10367*.
2. Levine et al. (2025). The Open Molecules 2025 (OMol25) Dataset, Evaluations, and Models. Meta FAIR. *arXiv:2505.08762*.
3. Yang et al. (2024). MatterSim: A Deep Learning Atomistic Model Across Elements, Temperatures, and Pressures. *arXiv:2405.04967*.
4. Anstine, Zubatyuk, Isayev (2024). AIMNet2: A Neural Network Potential to Meet Your Neutral, Charged, Organic, and Elemental-Organic Needs. *J. Chem. Inf. Model.* DOI: 10.1021/acs.jcim.4c00080.
5. Devereux et al. (2020). Extending the Applicability of the ANI Deep Learning Molecular Potential to Sulfur and Halogens. *J. Chem. Theory Comput.* 16, 4192–4202.
6. Galvelis et al. (2024). AceFF: A Neural Network Force Field for Protein-Ligand Systems. Acellera technical report.
7. Park et al. (2024). SevenNet-Omni: Scalable EQuivariant Neural Network Interatomic Potential. *Nature Communications* (in press 2026).  Seoul National University MDIL.
8. Grimme et al. (2017). A Robust and Accurate Tight-Binding Quantum Chemical Method. *JCTC* 13, 1989–2009 (GFN1-xTB).
9. Bannwarth et al. (2019). GFN2-xTB — An Accurate and Broadly Parametrized Self-Consistent Tight-Binding Quantum Chemical Method. *JCTC* 15, 1652–1671.
10. Spicher & Grimme (2020). Robust Atomistic Modeling of Materials, Organometallic, and Biochemical Systems. *Angew. Chem. Int. Ed.* 59, 15665–15673 (GFN-FF).
11. Halgren (1996). Merck Molecular Force Field. I. Basis, Form, Scope, Parameterization, and Performance of MMFF94. *J. Comput. Chem.* 17, 490–519.

12. Rappé et al. (1992). UFF, a Full Periodic Table Force Field for Molecular Mechanics and Molecular Dynamics Simulations. *JACS* 114, 10024–10035.
 13. Bursch, Mewes, Hansen, Grimme (2022). Best-Practice DFT Protocols for Basic Molecular Computational Chemistry. *Angew. Chem. Int. Ed.* 61, e202205735.
 14. Wohlfahrt-Mehrens, Becker, Mey et al. (2026). Boltz-2x: Physicality-Steered Cofold for Drug Discovery. (Boltz community release Mar 2026; preprint forthcoming.)
 15. Gaulton et al. (2017). The ChEMBL database in 2017. *Nucl. Acids Res.* 45, D945–D954.
 16. Adamo & Barone (1999). Toward Reliable Density Functional Methods Without Adjustable Parameters: The PBE0 Model. *J. Chem. Phys.* 110, 6158–6170. (Reference DFT methodology baseline.)
 17. Becke (1993). Density-Functional Thermochemistry. III. The Role of Exact Exchange. *J. Chem. Phys.* 98, 5648–5652. (B3LYP.)
 18. Grimme, Antony, Ehrlich, Krieg (2010). A consistent and accurate ab initio parametrization of density functional dispersion correction (DFT-D) for the 94 elements H-Pu. *J. Chem. Phys.* 132, 154104. (D3.)
 19. Grimme, Ehrlich, Goerigk (2011). Effect of the damping function in dispersion corrected density functional theory. *J. Comput. Chem.* 32, 1456–1465. (D3-BJ.)
-

Supplementary Material

S1. Full per-replicate per-pair Pearson tables (v11–v23)

(See companion `_metadata/paper_a_data_index.md` for paths to per-replicate CSV files. The headline subset of five pairs is given in Table 1 of the main text; the full $10 \times 10 \times 13$ -replicate tensor is in `pilot/round27_paperA/cluster_AB_analysis/v5g_13rep_full.npz`.)

S2. NNP version/build details

Method	Software	Version	Checkpoint weights	URL/DOI
GFN1-xTB	xtb	6.7.0	built-in	grimme-lab/xtb
GFN2-xTB	xtb	6.7.0	built-in	grimme-lab/xtb
GFN-FF	xtb	6.7.0	built-in	grimme-lab/xtb
AceFF-2	aceforce	2024.x	Acellera proprietary	acellera.com
AIMNet2-NSE	aimnet2	2024.x	aimnet2_b973c_d3_ens	github.com/ isayevlab/ AIMNet2
ANI-2x	torchani	2.2	ani2x ensemble	github.com/ aiqm/torchani
Orb-v3 OMat	orb-models	0.5.x	orb_v3_conservative_inf_github	github.com/orbital-materials/ orb-models
Orb-v3 OMol25	orb-models	0.5.x	orb_v3_conservative_inf_github	github.com/orbital-materials/ orb-models
MatterSim-5M	mattersim	1.1.x	MatterSim-v1.0.0-5M	github.com/microsoft/mattersim
SevenNet-OMat24	sevens	0.10.x	7net-omat-24	github.com/ MDIL-SNU/ SevenNet 
SevenNet-MPA	sevens	0.10.x	7net-mpa	as above
SevenNet-Mat-PES-PBE	sevens	0.10.x	7net-matpes-pbe	as above
SevenNet-OMol25-high	sevens	0.10.x	7net-omol25-high	as above
FeNNix-Bio1S	fennix	0.3.x	Bio1S	github.com/sorbonne-cnrs/fennix
MMFF94	RDKit	2024.09	MMFFGetMoleculeForceField	rdkit.org
UFF	RDKit	2024.09	UFFGetMoleculeForceField	rdkit.org

S3. Bug-fix and gotcha notes (for reproducibility)

- **RTX 5090 sm_120 / torch 2.6 kernel mismatch.** Orb-v3 and SevenNet require torch $\geq$ 2.7.0 cu128 wheels (or TORCHDYNAMO_DISABLE=1 and CPU fallback) on Blackwell GPUs. See `feedback_rtx5090_sm120_torch_kernel_compat`.
- **Orb-v3 OMol25 silent fail on missing charge/spin.** Atoms must have `atoms.info["charge"]` and `atoms.info["spin"]` populated; otherwise OMol25 silently runs with `charge=0/spin=0`, which is wrong for ionized hydroxamates and for Zn-containing extracts. See `feedback_orb_v3_omol_charge_spin`.
- **MatterSim per-atom spin requirement.** MatterSim accepts unannotated atoms; no special handling needed for our Zn^{2+} (d^{10} closed-shell).
- **GFN2-xTB SCF divergence on full protein+ligand complex.** ~3000-atom protein+ligand complexes fail xTB GFN1/GFN2 with SIGSEGV (SCF divergence). GFN-FF works. For all NNP single-point evaluations we use the ligand-only extract (heavy atoms + Zn^{2+} if within 6 Å). See `feedback_xtb_gfn2_protein_complex_segfault`.
- **GFN-FF disk explosion.** xTB GFN-FF writes a ~93 MB `gfnff_topo` file per PDB and does not auto-delete; cleanup must be scripted. See `feedback_xtb_gfnff_disk_explosion`.
- **Across-replicate variance is small but not zero.** SD across 13 Boltz seeds is 0.001–0.038 on all pairs; the cluster topology is stable.

S4. Replicate-to-Boltz-seed mapping

Replicate ID	Boltz seed
v11	42
v12	43
v13	44
v14	45
v15	46
v16	47
v17	48
v18	49
v19	50
v20	51
v21	52
v22	53
v23	54

S5. DFT reference data (subset)

CSV: `/home/crazat/genesis_medicine/pilot/round27_paperA/dft_b3lyp_d3_verdict.csv`.

Headline rows for ChEMBL406 v22:

Con- former	OMat (eV)	OMol25 (eV)	GFN1 (Eh)	B3LYP/def2- SVP (Ha)	B3LYP- D3(BJ)/def2- SVP (Ha)
model_98	-213.373	-47856.230	-80.310	-1785.6547	(not run)
model_36	-213.301	-47856.246	-80.303	-1785.6804	-1785.8444
model_82	-213.842	-47858.637	-80.338	-1785.6600	-1785.8256
model_65	-212.954	-47858.066	-80.288	-1785.6814	-1785.8445
model_79	-212.673	-47857.559	-80.274	-1785.6964	-1785.8597

Pearson r against B3LYP/def2-SVP on ChEMBL406 (n = 5): OMat = -0.84, GFN1-xTB = -0.84, OMol25 = +0.09 (all marginal at $p \approx 0.07$; cluster claim does not rest on this).

S6. Companion files and analysis pipelines

See `_metadata/paper_a_data_index.md` for a structured index of CSV outputs, analysis scripts, and figure-generation scripts.

Manuscript draft v0.1, 2026-05-12. Author: Han, Cheongwoo. Comments to crazat7@gmail.com.